

Heart Disease Classification

1. Problem Definition: Create ML model to predict if a person has heart disease
2. Data: From kaggle ("<https://www.kaggle.com/datasets/kamilpytlak/personal-key-indicators-of-heart-disease>")
3. Evaluation: Predict with 90% accuracy
4. Features: BMI, Smoking, Alcohol Drinking, Stroke, Physical Health, Mental Health, Diff Walking, Sex, Age Category, Race, Diabetic, Physical Activity, Gen Health, Sleep Time, Asthma, Kidney Disease, Skin Cancer

```
In [1]: # Importing tools
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.utils.class_weight import compute_class_weight
from imblearn.under_sampling import RandomUnderSampler
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
import pickle
```

```
In [2]: # Loading data into dataframe
heart_df = pd.read_csv("heart_disease_data.csv")
heart_df.head(5)
```

```
Out[2]:
```

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth
0	No	16.60	Yes	No	No	3.0	3.0
1	No	20.34	No	No	Yes	0.0	0.0
2	No	26.58	Yes	No	No	20.0	3.0
3	No	24.21	No	No	No	0.0	0.0
4	No	23.71	No	No	No	28.0	0.0

Exploring the data

```
In [3]: # Check amount of data
len(heart_df)
```

```
Out[3]: 319795
```

```
In [4]: heart_df.describe()
```

```
Out[4]:
```

	BMI	PhysicalHealth	MentalHealth	SleepTime
count	319795.000000	319795.000000	319795.000000	319795.000000
mean	28.325399	3.37171	3.898366	7.097075
std	6.356100	7.95085	7.955235	1.436007
min	12.020000	0.00000	0.000000	1.000000
25%	24.030000	0.00000	0.000000	6.000000
50%	27.340000	0.00000	0.000000	7.000000
75%	31.420000	2.00000	3.000000	8.000000
max	94.850000	30.00000	30.000000	24.000000

```
In [5]: # Check for missing data
heart_df.isna().sum()
```

```
Out[5]: HeartDisease      0
BMI                      0
Smoking                  0
AlcoholDrinking          0
Stroke                   0
PhysicalHealth            0
MentalHealth              0
DiffWalking              0
Sex                      0
AgeCategory              0
Race                     0
Diabetic                 0
PhysicalActivity          0
GenHealth                0
SleepTime                0
Asthma                   0
KidneyDisease            0
SkinCancer               0
dtype: int64
```

```
In [6]: # Check for biases in data
heart_df["HeartDisease"].value_counts()
```

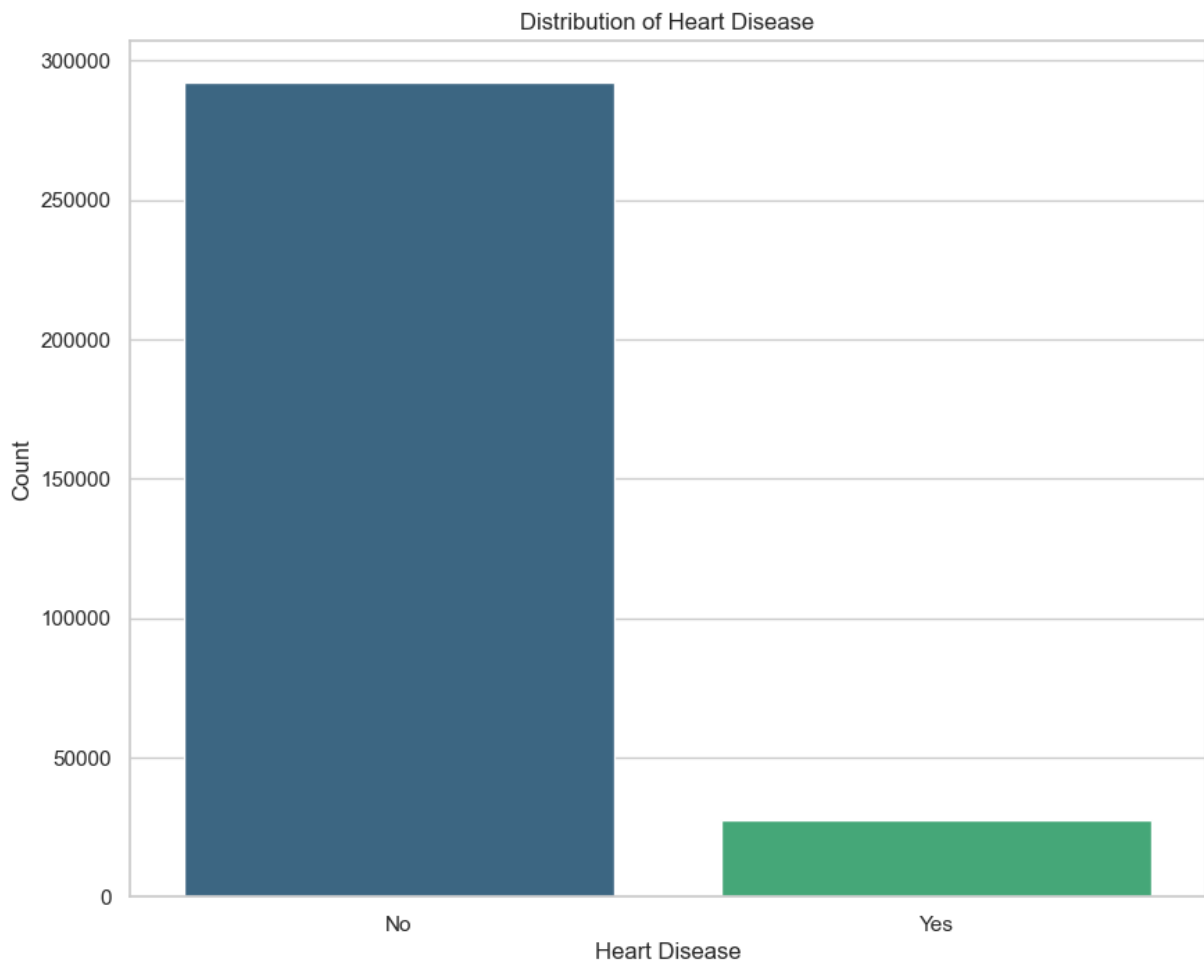
```
Out[6]: HeartDisease
No      292422
Yes     27373
Name: count, dtype: int64
```

There is a significant amount of "no" heart disease than "yes", so the training data should be randomly under-sampled

Graphs

```
In [7]: # Set the style for seaborn
sns.set(style="whitegrid")

# Distribution of Heart Disease
plt.figure(figsize = (10, 8))
sns.countplot(x = 'HeartDisease', data = heart_df, hue = 'HeartDisease', palette = 'magma')
plt.title('Distribution of Heart Disease')
plt.xlabel('Heart Disease')
plt.ylabel('Count')
plt.show()
```



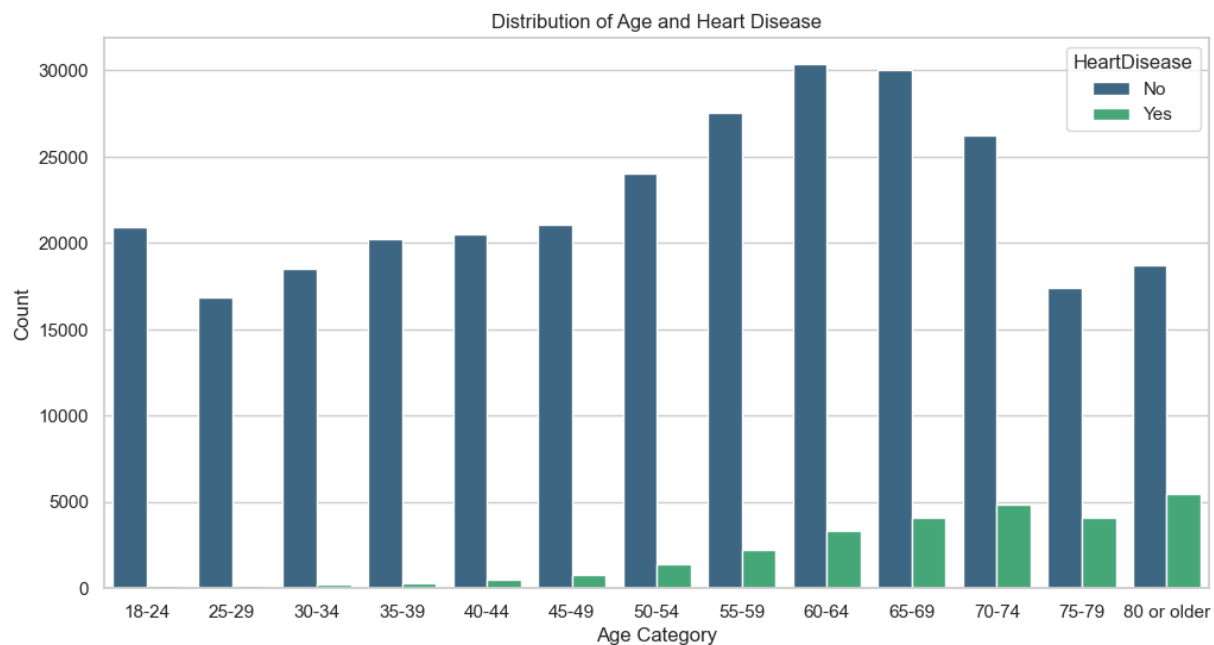
```
In [8]: plt.figure(figsize = (12, 6))

# sort by age
age_order = sorted(heart_df['AgeCategory'].unique())

sns.countplot(x = 'AgeCategory', hue = 'HeartDisease', data = heart_df, order = age_order)

# title and labels
plt.title('Distribution of Age and Heart Disease')
plt.xlabel('Age Category')
```

```
plt.ylabel('Count')
plt.show()
```



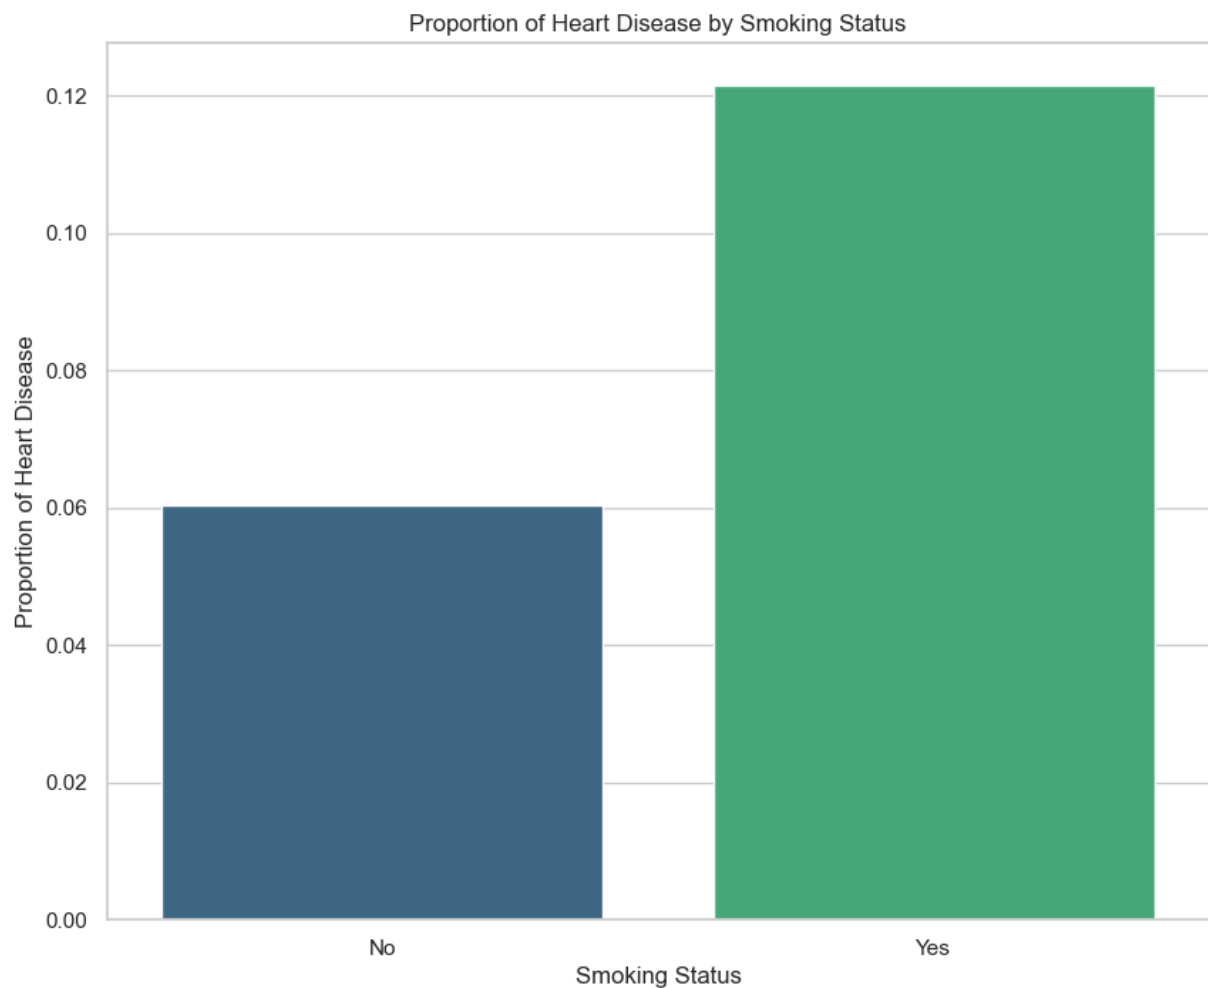
```
In [9]: plt.figure(figsize = (10, 8))

# Calculate the proportion of heart disease within each smoking category
heart_disease_proportion = heart_df.groupby('Smoking')['HeartDisease'].value

# Creates bar graph
sns.barplot(x = heart_disease_proportion.index, y = heart_disease_proportion)

# Title and labels
plt.title('Proportion of Heart Disease by Smoking Status')
plt.xlabel('Smoking Status')
plt.ylabel('Proportion of Heart Disease')

plt.show()
```



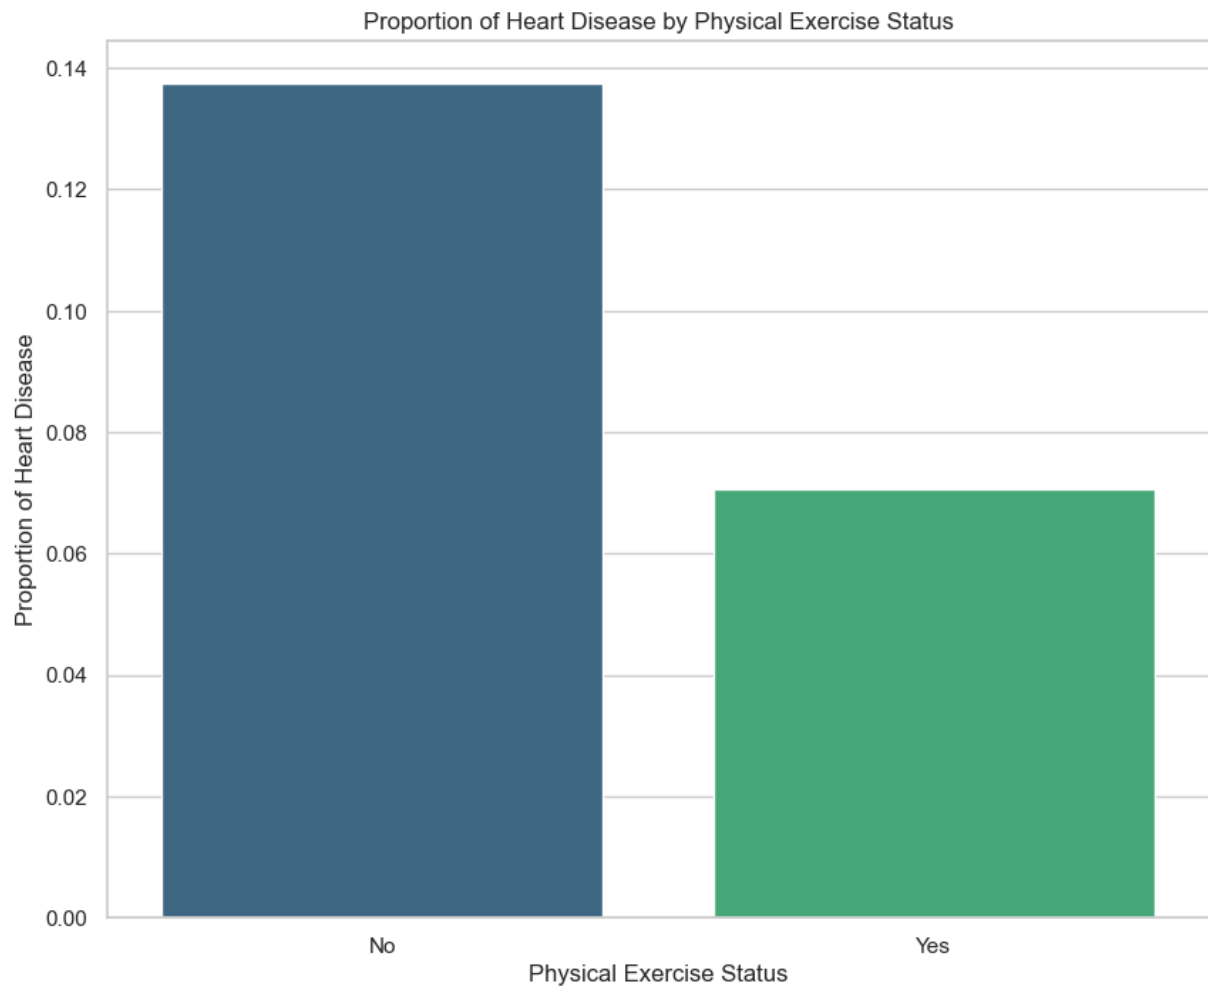
```
In [10]: plt.figure(figsize = (10, 8))

# Calculate proportion of heart disease based on physical activity status
heart_disease_proportion_exercise = heart_df.groupby('PhysicalActivity')['HeartDisease'].agg('mean')

sns.barplot(x = heart_disease_proportion_exercise.index, y = heart_disease_proportion_exercise.values)

plt.title('Proportion of Heart Disease by Physical Exercise Status')
plt.xlabel('Physical Exercise Status')
plt.ylabel('Proportion of Heart Disease')

plt.show()
```



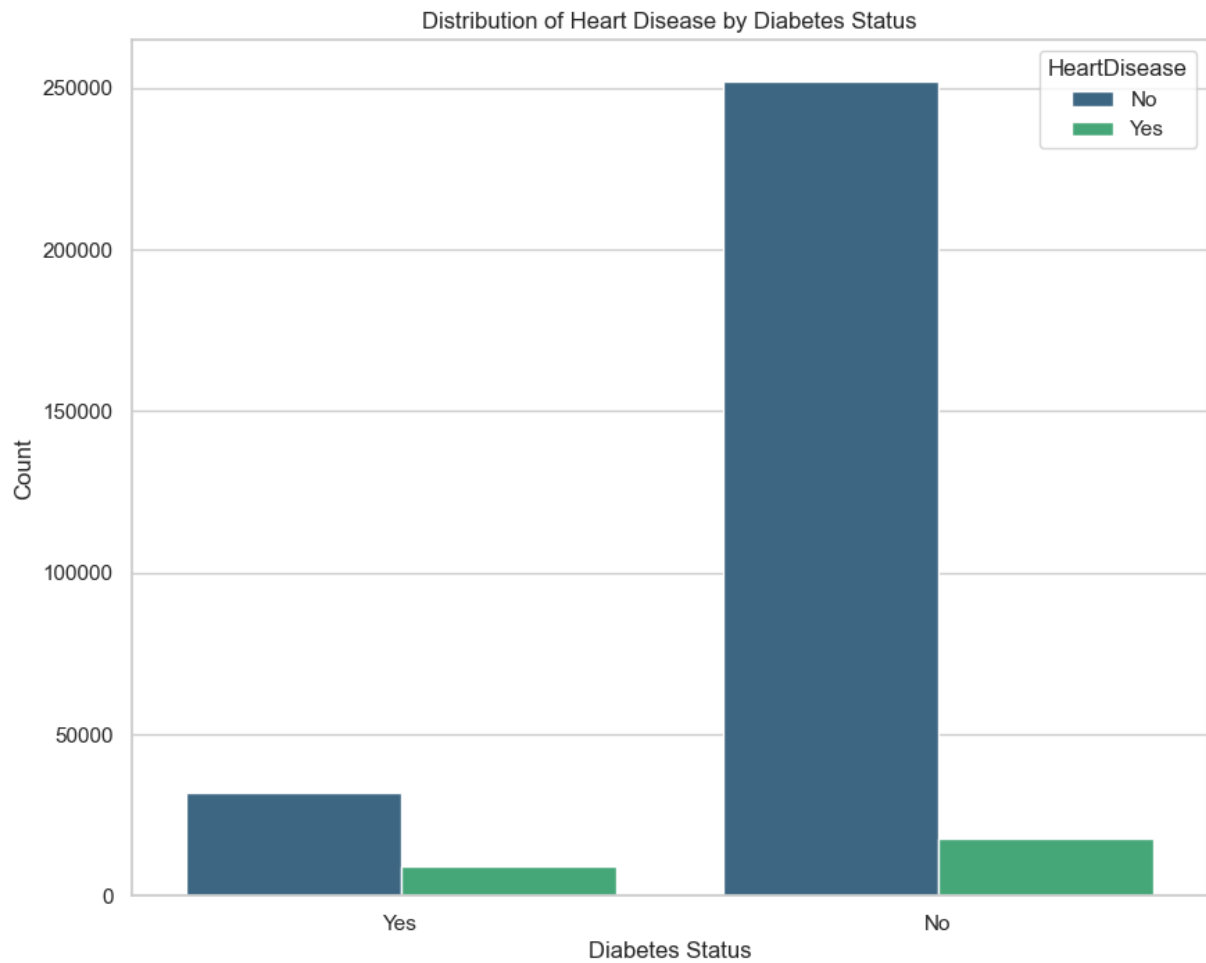
```
In [11]: # Filter the DataFrame for 'Yes' and 'No' in the 'Diabetic' column
filtered_heart_df = heart_df[heart_df['Diabetic'].isin(['Yes', 'No'])]

plt.figure(figsize = (10, 8))

sns.countplot(x = 'Diabetic', hue = 'HeartDisease', data = filtered_heart_df)

plt.title('Distribution of Heart Disease by Diabetes Status')
plt.xlabel('Diabetes Status')
plt.ylabel('Count')

plt.show()
```

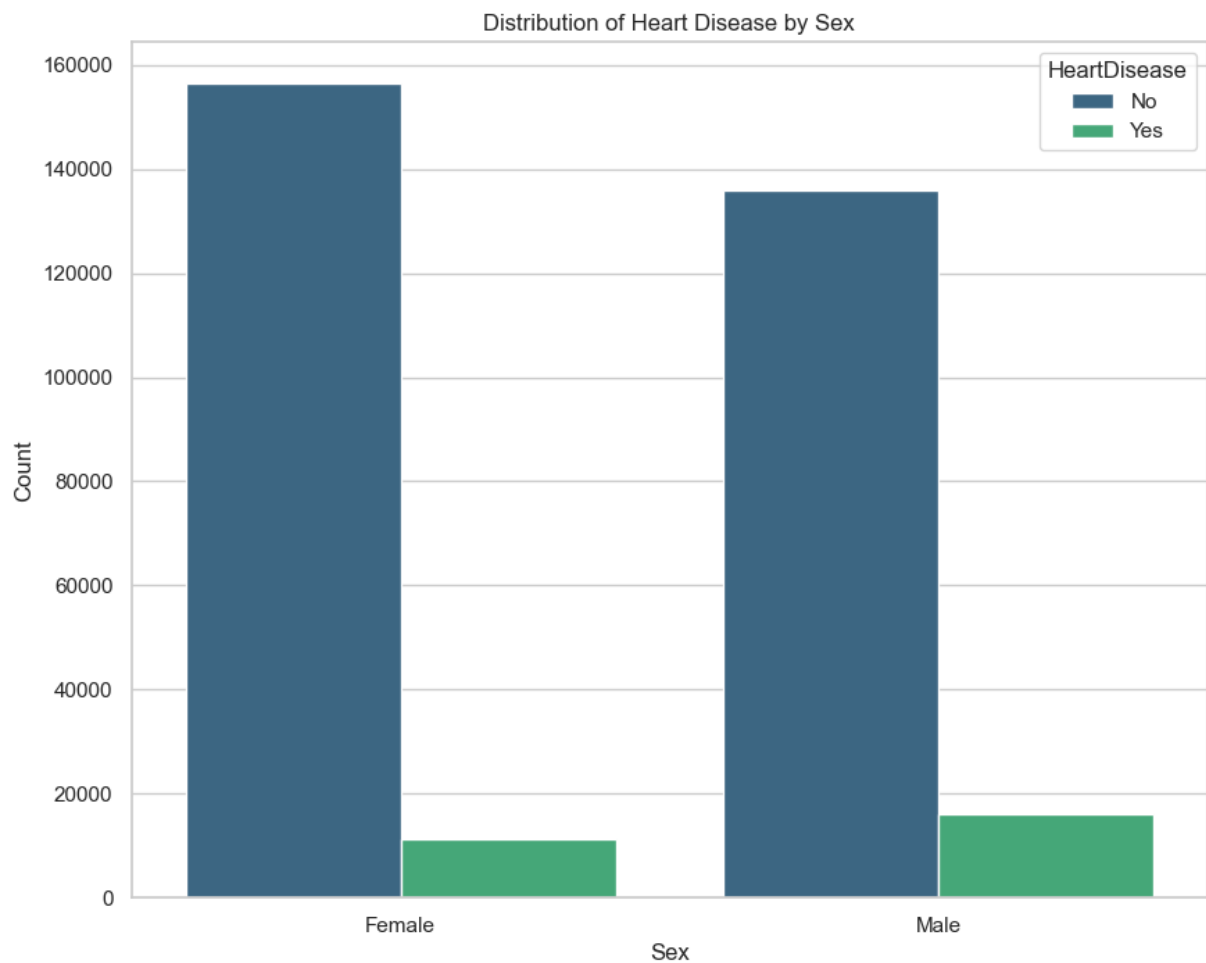


```
In [12]: plt.figure(figsize = (10, 8))

sns.countplot(x = 'Sex', hue = 'HeartDisease', data = heart_df, palette = "v

plt.title('Distribution of Heart Disease by Sex')
plt.xlabel('Sex')
plt.ylabel('Count')

plt.show()
```



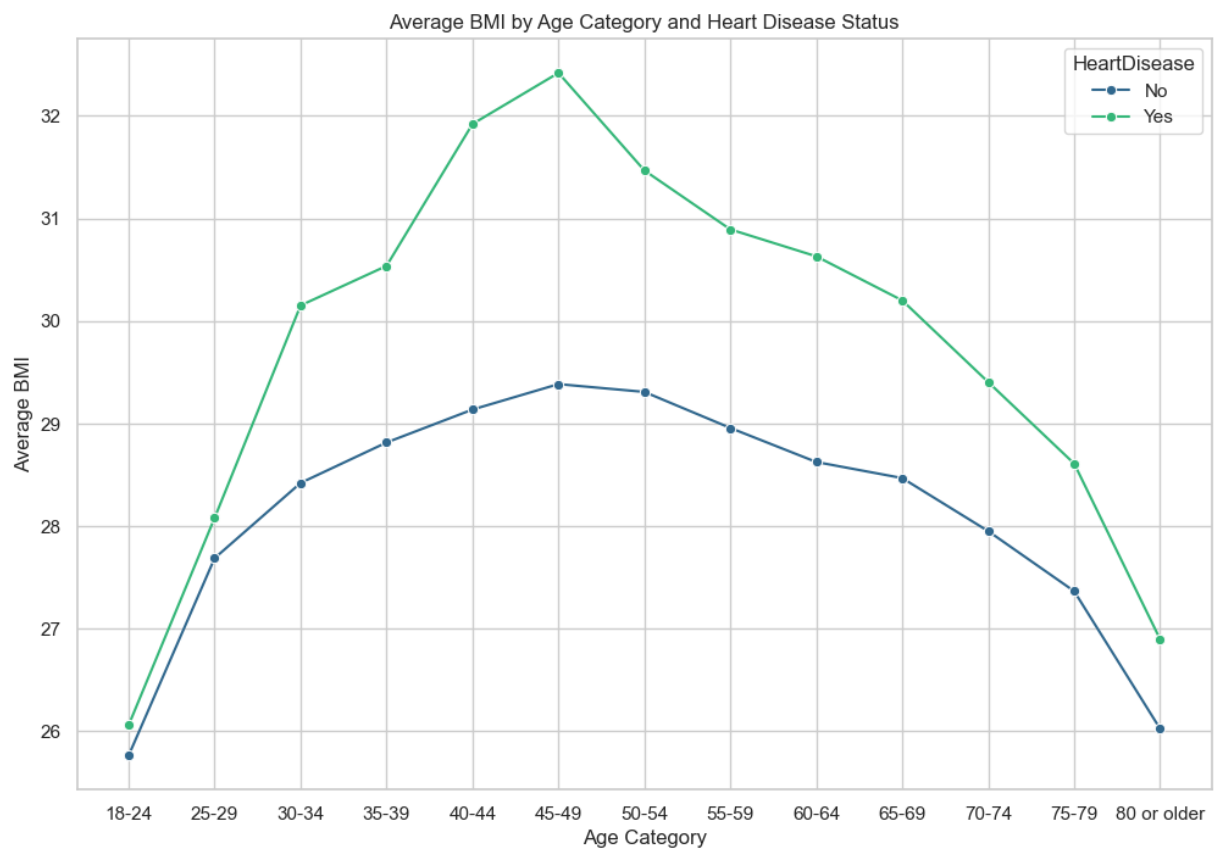
```
In [13]: plt.figure(figsize = (12, 8))

# Group the data by heart disease status and age category and calculate the
grouped_data_bmi = heart_df.groupby(['HeartDisease', 'AgeCategory'])['BMI'].

# Line plot
sns.lineplot(x = 'AgeCategory', y = 'BMI', hue = 'HeartDisease', data = grou

plt.title('Average BMI by Age Category and Heart Disease Status')
plt.xlabel('Age Category')
plt.ylabel('Average BMI')

plt.show()
```

Converting data to numbers

In [14]: `heart_df.head()`

Out[14]:

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth
0	No	16.60	Yes	No	No	3.0	3.0
1	No	20.34	No	No	Yes	0.0	0.0
2	No	26.58	Yes	No	No	20.0	3.0
3	No	24.21	No	No	No	0.0	0.0
4	No	23.71	No	No	No	28.0	0.0

```
In [15]: bool_to_num = {'Yes': 1, 'No': 0}

for column in heart_df.columns:
    if set(heart_df[column].unique()) == {'Yes', 'No'}:
        heart_df[column] = heart_df[column].map(bool_to_num)

heart_df.head()
```

Out [15]:

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHea
0	0	16.60	1	0	0	3.0	30
1	0	20.34	0	0	1	0.0	0
2	0	26.58	1	0	0	20.0	30
3	0	24.21	0	0	0	0.0	0
4	0	23.71	0	0	0	28.0	0

In [16]:

```
# Instantiate label encoder
label_encoder = LabelEncoder()

columns_to_encode = ['Sex', 'AgeCategory', 'GenHealth', 'Race', 'Diabetic']

# Loop through all columns to encode, then encode
for column in columns_to_encode:
    heart_df[column] = label_encoder.fit_transform(heart_df[column])

heart_df.head()
```

Out [16]:

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHea
0	0	16.60	1	0	0	3.0	30
1	0	20.34	0	0	1	0.0	0
2	0	26.58	1	0	0	20.0	30
3	0	24.21	0	0	0	0.0	0
4	0	23.71	0	0	0	28.0	0

Modeling

Using Random Forest regression algorithm to handle imbalanced data

In [17]:

```
# Random Forest
X = heart_df.drop("HeartDisease", axis = 1)
y = heart_df["HeartDisease"]

# Split 80% train, 20% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, r

rf_clf = RandomForestClassifier(n_estimators = 10)
rf_clf.fit(X_train, y_train)

rf_clf.score(X_test, y_test)
```

Out [17]: 0.9049078315796056

Evaluation

```
In [18]: # Predicted outputs
y_pred = rf_clf.predict(X_test)

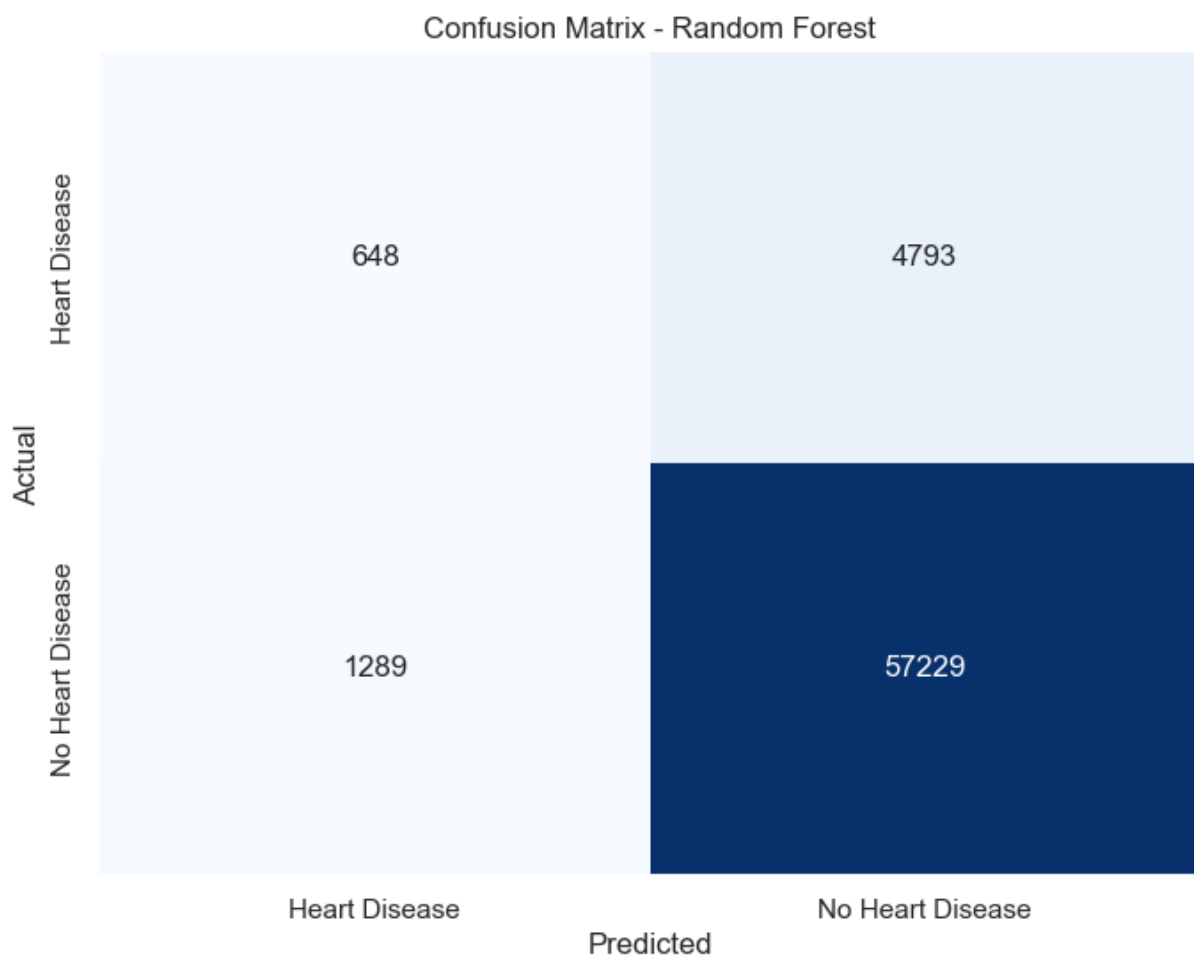
cm = confusion_matrix(y_test, y_pred)

# Flip the confusion matrix
cm_flipped = cm[::-1, ::-1]

plt.figure(figsize = (8, 6))

sns.heatmap(cm_flipped, annot = True, fmt = "d", cmap = "Blues", cbar = False,
            xticklabels = ['Heart Disease', 'No Heart Disease'],
            yticklabels = ['Heart Disease', 'No Heart Disease'])

plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



Due to the dataset's predominant instances without heart disease, we observe an increased number of false negatives.

```
In [19]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	58518
1	0.33	0.12	0.18	5441
accuracy			0.90	63959
macro avg	0.63	0.55	0.56	63959
weighted avg	0.87	0.90	0.88	63959

Experimentation

Class Weights, assigns higher weight to "yes heart disease"

```
In [20]: rf_clf_cw = RandomForestClassifier(n_estimators = 10, class_weight = {0: 1,
rf_clf_cw.fit(X_train, y_train)

rf_clf_cw.score(X_test, y_test)
```

```
Out[20]: 0.9043606060132272
```

```
In [21]: rf_clf_cw.score(X_test, y_test) > rf_clf.score(X_test, y_test)
```

```
Out[21]: False
```

Random undersampling "no heart disease"

```
In [22]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Apply random undersampling to the training set
rus = RandomUnderSampler(random_state = 24)
X_resampled, y_resampled = rus.fit_resample(X_train, y_train)

# Train Random Forest on resampled data
rf_clf_rs = RandomForestClassifier(n_estimators = 10)
rf_clf_rs.fit(X_resampled, y_resampled)

rf_clf_rs.score(X_test, y_test)
```

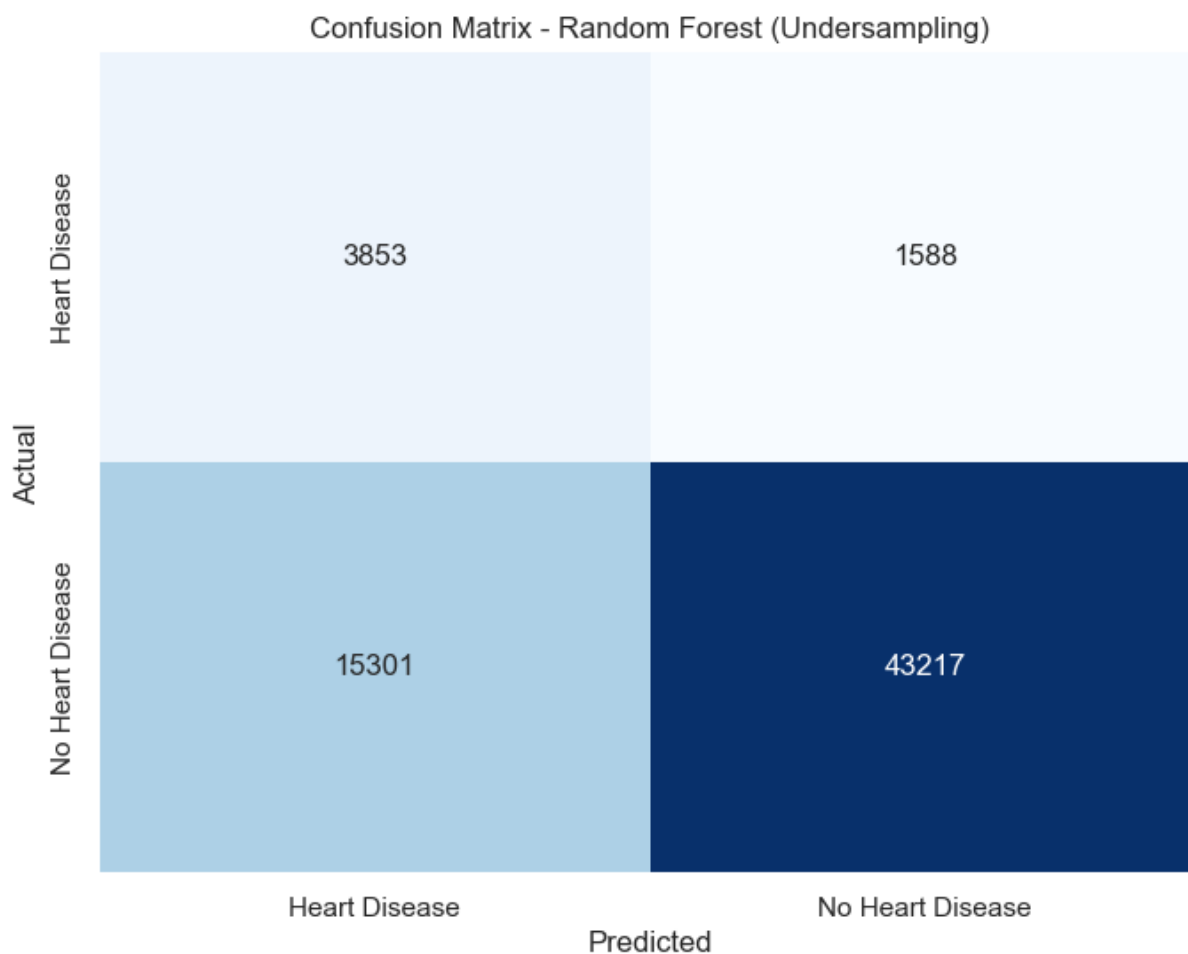
```
Out[22]: 0.7359402116981191
```

```
In [23]: y_pred_rs = rf_clf_rs.predict(X_test)

cm = confusion_matrix(y_test, y_pred_rs)

plt.figure(figsize=(8, 6))
sns.heatmap(cm[::-1, ::-1], annot = True, fmt = "d", cmap = "Blues", cbar =
            xticklabels = ['Heart Disease', 'No Heart Disease'],
            yticklabels = ['Heart Disease', 'No Heart Disease'])
plt.title('Confusion Matrix - Random Forest (Undersampling)')
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
plt.show()
```



```
In [24]: print(classification_report(y_test, y_pred_rs))
```

	precision	recall	f1-score	support
0	0.96	0.74	0.84	58518
1	0.20	0.71	0.31	5441
accuracy			0.74	63959
macro avg	0.58	0.72	0.57	63959
weighted avg	0.90	0.74	0.79	63959

Tuning hyper-parameters

```
In [25]: # Try 100-200 in increments of 10 for n-estimators
for i in range(100, 200, 10):
    print(f"Trying model with {i} estimators...")
    model = RandomForestClassifier(n_estimators = i, random_state = 24).fit(
    print(f"Model accuracy on test set: {model.score(X_test, y_test) * 100:.2f}%")
    print("")
```


Out[28]: 0.9155709126158945

```
In [29]: # Trying a more extensive grid search
param_grid = {'n_estimators': [20, 50, 170, 200], 'min_samples_split': [3, 4]

grid_search = GridSearchCV(estimator = rf_clf, param_grid = param_grid)
grid_search.fit(X_train, y_train)

grid_search.best_params_
```

Out[29]: {'min_samples_leaf': 5, 'min_samples_split': 3, 'n_estimators': 170}

```
In [30]: # Best estimators from grid search
rf_clf = RandomForestClassifier(n_estimators = 50, min_samples_leaf = 5, min
rf_clf.fit(X_train, y_train)

rf_clf.score(X_test, y_test)
```

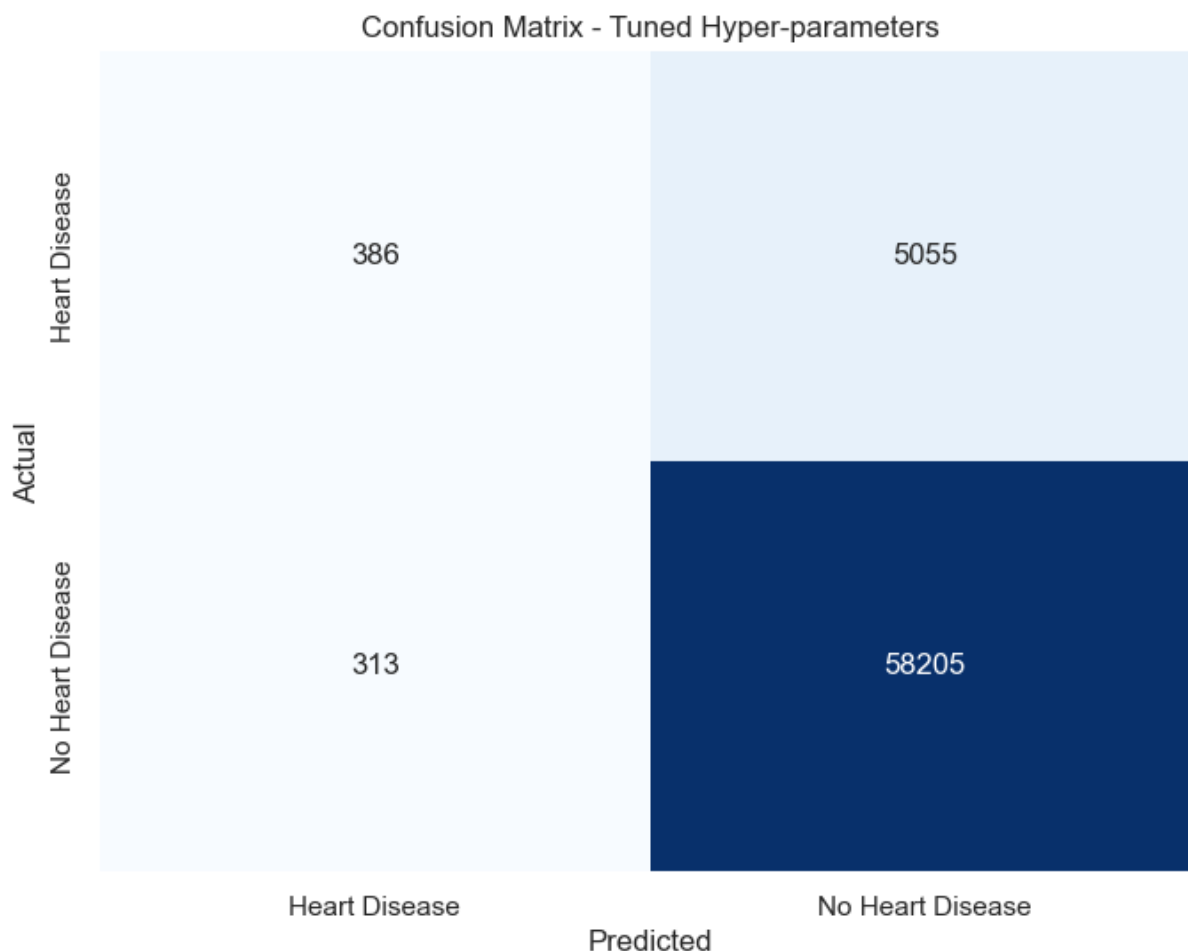
Out[30]: 0.9160712331337263

```
In [31]: y_pred = rf_clf.predict(X_test)

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))
sns.heatmap(cm[::-1, ::-1], annot = True, fmt = "d", cmap = "Blues", cbar =
            xticklabels = ['Heart Disease', 'No Heart Disease'],
            yticklabels = ['Heart Disease', 'No Heart Disease'])
plt.title('Confusion Matrix - Tuned Hyper-parameters')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0	0.92	0.99	0.96	58518
1	0.55	0.07	0.13	5441
accuracy			0.92	63959
macro avg	0.74	0.53	0.54	63959
weighted avg	0.89	0.92	0.89	63959

```
In [32]: # Parameters to optimize
param_dist = {
    'n_estimators': np.random.randint(10, 100, size = 10),
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': np.random.randint(2, 100, size = 13),
    'min_samples_leaf': np.random.randint(1, 100, size = 13),
    'class_weight': [{0: 1, 1: w} for w in [5, 15, 20, 30, 50]]
}

rf_clf = RandomForestClassifier(random_state = 24)

random_search = RandomizedSearchCV(rf_clf, param_distributions = param_dist,
random_search.fit(X_train, y_train)

print("Best Parameters:", random_search.best_params_)

best_rf_clf_random = random_search.best_estimator_
```



```
best_rf_clf_random_score = best_rf_clf_random.score(X_test, y_test)
print("Best Model Score:", best_rf_clf_random_score)

y_pred_best_random = best_rf_clf_random.predict(X_test)
```

Best Parameters: {'n_estimators': 66, 'min_samples_split': 25, 'min_samples_leaf': 41, 'max_depth': 40, 'class_weight': {0: 1, 1: 5}}

Best Model Score: 0.8509982957832362

Best Hyperparameters

After conducting a grid search, the following hyperparameters were identified as the best for the Random Forest model:

- `n_estimators`: 50
- `min_samples_split`: 5
- `min_samples_leaf`: 6

Data Downsampling: Balancing Heart Disease Instances

```
In [33]: np.random.seed(24)
heart_disease_index = heart_df[heart_df['HeartDisease'] == 0].index

# Randomly select 150000 thousand rows
random_rows_remove = np.random.choice(heart_disease_index, size = 150000, replace = True)

# Remove the selected rows
heart_df_down_scaled = heart_df.drop(index = random_rows_remove)

# Check Shape and counts
print(f"Original DataFrame shape: {heart_df.shape}")
print(f"Updated DataFrame shape: {heart_df_down_scaled.shape}")
print("\nThe distribution of 'HeartDisease' in the original DataFrame:")
print(heart_df["HeartDisease"].value_counts())
print("\nThe distribution of 'HeartDisease' in the updated DataFrame:")
print(heart_df_down_scaled["HeartDisease"].value_counts())
```

Original DataFrame shape: (319795, 18)

Updated DataFrame shape: (169795, 18)

The distribution of 'HeartDisease' in the original DataFrame:

HeartDisease

0 292422

1 27373

Name: count, dtype: int64

The distribution of 'HeartDisease' in the updated DataFrame:

HeartDisease

0 142422

1 27373

Name: count, dtype: int64

```
In [34]: # Split down scaled data
X_rs = heart_df_down_scaled.drop("HeartDisease", axis = 1)
y_rs = heart_df_down_scaled["HeartDisease"]

X_train, X_test, y_train, y_test = train_test_split(X_rs, y_rs, test_size =

# Train new data set
rf_clf = RandomForestClassifier(n_estimators = 50, min_samples_leaf = 5, min
rf_clf.fit(X_train, y_train)

rf_clf.score(X_test, y_test)
```

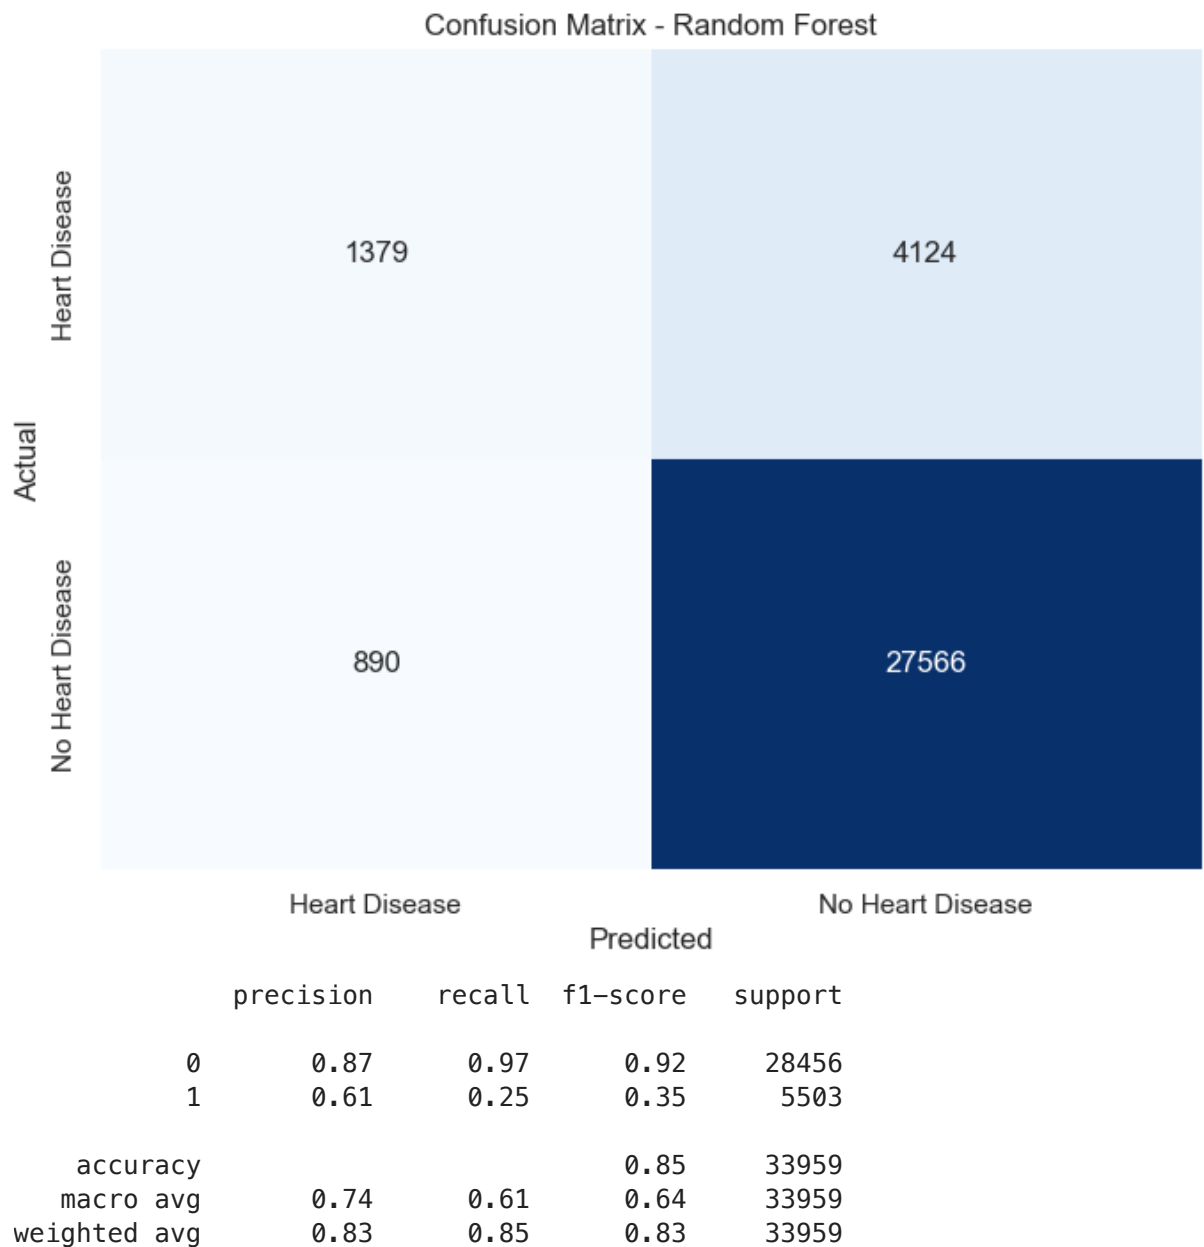
Out[34]: 0.8523513648811802

```
In [35]: # Predictions on the test set
y_preds = rf_clf.predict(X_test)

# Generate confusion matrix
cm = confusion_matrix(y_test, y_preds)

plt.figure(figsize=(8, 6))
sns.heatmap(cm[::-1, ::-1], annot = True, fmt = "d", cmap = "Blues", cbar =
            xticklabels = ['Heart Disease', 'No Heart Disease'],
            yticklabels = ['Heart Disease', 'No Heart Disease'])
plt.title('Confusion Matrix - Random Forest')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

print(classification_report(y_test, y_preds))
```



```
In [36]: # Cross validation, split in 5
cv_scores = cross_val_score(rf_clf, X, y, cv = 5, scoring = 'accuracy')

print("Cross-Validation Scores:", cv_scores)
print("Mean Cross-Validation Score:", cv_scores.mean())
print("")

cv_scores_r = cross_val_score(rf_clf, X_rs, y_rs, cv = 5, scoring = 'accuracy')

print("Cross-Validation Scores (reduced data):", cv_scores_r)
print("Mean Cross-Validation Score (reduced data):", cv_scores_r.mean())
```

Cross-Validation Scores: [0.91502369 0.91525821 0.91525821 0.91546147 0.91557091]

Mean Cross-Validation Score: 0.9153144983505058

Cross-Validation Scores (reduced data): [0.85296976 0.85061398 0.85276363 0.85135016 0.85167408]

Mean Cross-Validation Score (reduced data): 0.8518743190317736

Summary

1. Key Findings:

- The best-trained model achieved an accuracy of approximately 91.6%, providing a reliable predictive capability for heart disease.
- The precision-recall-f1 score evaluation suggests that while the model performs exceptionally well in identifying cases without heart disease (Class 0), it faces challenges in correctly classifying positive cases (Class 1).
- Smoking is identified as a significant factor contributing to heart disease, as it causes the proportion of heart disease cases to double.
- Engaging in physical activity appears to have a protective effect, reducing the proportion of heart disease cases by a noticeable amount.
- Individuals with a higher BMI exhibit a higher likelihood of heart disease, indicating that BMI is positively correlated with the risk of developing heart-related issues.

2. Insights:

- The analysis reveals a stark contrast in heart disease risk between smokers and non-smokers, emphasizing the importance of lifestyle choices.
- Physical activity emerges as a crucial factor in mitigating heart disease risk, reinforcing the significance of regular exercise.
- The correlation between BMI and heart disease highlights the need for weight management as a preventive measure.

3. Limitations:

- The dataset exhibits a notable imbalance, with a significantly larger number of samples labeled as "No heart disease" compared to "Yes heart disease." This imbalance could impact the model's performance, particularly in detecting positive cases.

4. Next Steps:

- Address data bias by collecting additional samples for individuals with heart disease to create a more balanced dataset.
- Explore advanced techniques, such as oversampling or undersampling, to mitigate the impact of data imbalance on model training.

5. Acknowledgments:

- The dataset used in this analysis was obtained from Kaggle, with appreciation for the contribution of Kamil Pytlak. The dataset link is [here](#).

This comprehensive summary provides a clear overview of the key findings, insights, recommendations, limitations, next steps, and acknowledgments derived from the analysis of personal key indicators of heart disease.